

Lecture 31: Concurrency and Parallelism (I)

Announcement: HKN (Eta Kappa Nu) will be holding a review session for the upcoming midterm this Sunday, 4/8, from 1100-1300 in 306 Soda Hall (HP Auditorium). Bring yourselves, some pencil and paper, and any questions that you may have.

Definitions

- *Sequential Process*: Our subject matter up to now: processes that (ultimately) proceed in *a single sequence* of primitive steps.
- *Concurrent Processing*: The logical or physical division of a process into multiple sequential processes.
- *Parallel Processing*: A variety of concurrent processing characterized by the *simultaneous execution* of sequential processes.
- *Distributed Processing*: A variety of concurrent processing in which the individual processes are physically separated (often using heterogeneous platforms) and communicate through some network structure.

Purposes

We may divide a single program into multiple programs for various reasons:

- *Computation Speed* through operating on separate parts of a problem simultaneously, or through
- *Communication Speed* through putting parts of a computation near the various data they use.
- *Reliability* through having multiple physical copies of processing or data.
- *Security* through separating sensitive data from untrustworthy users or processors of data.
- *Better Program Structure* through decomposition of a program into logically separate processes.
- *Resource Sharing* through separation of a component that can serve multiple users.
- *Manageability* through separation (and sharing) of components that may need frequent updates or complex configuration.


Communicating Sequential Processes

- All forms of concurrent computation can be considered instances of *communicating sequential processes*.
- That is, a bunch of “ordinary” programs that communicate with each other through what is, from their point of view, input and output operations.
- Sometimes the actual communication medium is *shared memory*: input looks like reading a variable and output looks like writing a variable. In both cases, the variable is in memory accessed by multiple computers.
- At other times, communication can involve I/O over a network such as the Internet.
- In principle, either underlying mechanism can be made to look like either access to variables or explicit I/O operations to a programmer.

Distributed Communication

- With sequential programming, we don't think much about the cost of "communicating" with a variable; it happens at some fixed speed that is (we hope) related to the processing speed of our system.
- With distributed computing, the architecture of communication becomes important.
- In particular, costs can become uncertain or heterogeneous:
 - It may take longer for one pair of components to communicate than for another, or
 - The communication time may be unpredictable or load-dependent.

Simple Client-Server Models



- Example: web servers
- Good for providing a service
- Many clients, one server
- Easy server maintenance.
- Single point of failure
- Problems with scaling

Variations: on to the cloud

- Google and other providers modify this model with redundancy in many ways.
- For example, *DNS load balancing* (DNS = *Domain Name System*) allows us to specify multiple servers.
- Requests from clients go to different servers that all have copies of relevant information.
- Put enough servers in one place, you have a *server farm*. Put servers in lots of places, and we have a *cloud*.

Communication Protocols

- One characteristic of modern distributed systems is that they are conglomerations of products from many sources.
- Web browsers are a kind of universal client, but there are numerous kinds of browsers and many potential servers (and clouds of servers).
- So there must be some agreement on how they talk to each other.
- The *IP Protocol* is an agreement for specifying destinations, packaging messages, and delivering those messages.
- On top of this, the *transmission control protocol (TCP)* handles issues like persistent telephone-like connections and congestion control.
- The DNS handles conversions between names (inst.eecs.berkeley.edu) and IP addresses (128.32.42.199).
- The *HyperText Transfer Protocol* handles transfer of requests and responses from web servers.

Example: HTTP

- When you click on a link, such as

`http://inst.eecs.berkeley.edu/~cs61a/lectures,`

your browser:

- Consults the DNS to find out where to look for `inst.eecs.berkeley.edu`.
- Sends a message to port 80 at that address:

`GET ~cs61a/lectures HTTP 1.1`

- The program listening there (the web server) then responds with

`HTTP/1.1 200 OK`


`Content-Type: text/html`

`Content-Length: 1354`

`<html> ... text of web page`


- Protocol has other messages: for example, POST is often used to send data in forms from your browser. The data follows the POST message and other headers.

Peer-to-Peer Communication



- No central point of failure; clients talk to each other.
- Can route around network failures.
- Computation and memory shared.
- Can grow or shrink as needed.
- Used for file-sharing applications, bot-nets (!).
- But, deciding routes, avoiding congestion, can be tricky.
- (E.g., Simple scheme, broadcasting all communications to everyone, requires N^2 communication resource. Not practical.
- Maintaining consistency of copies requires work.
- Security issues.

Clustering



- A peer-to-peer network of “supernodes,” each serving as a server for a bunch of clients.
- Allows scaling; could be nested to more levels.
- Examples: Skype, network time service.