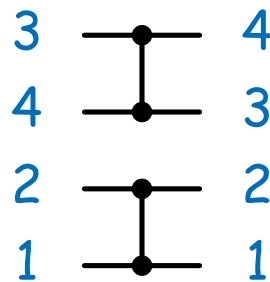


## Lecture 32: Parallelism

- Moore's law ("Transistors per chip doubles every  $N$  years"), where  $N$  is roughly 2 (about  $1,000,000\times$  increase since 1971).
- Has also applied to processor speeds (doubling a bit faster).
- But predicted to flatten: further increases to be obtained through *parallel processing* (witness: multicore/manycore processors).
- With distributed processing, issues involve interfaces, reliability, communication issues.
- With other parallel computing, where the aim is performance, issues involve synchronization, balancing loads among processors, and, yes, "data choreography" and communication costs.

# Example of Parallelism: Sorting

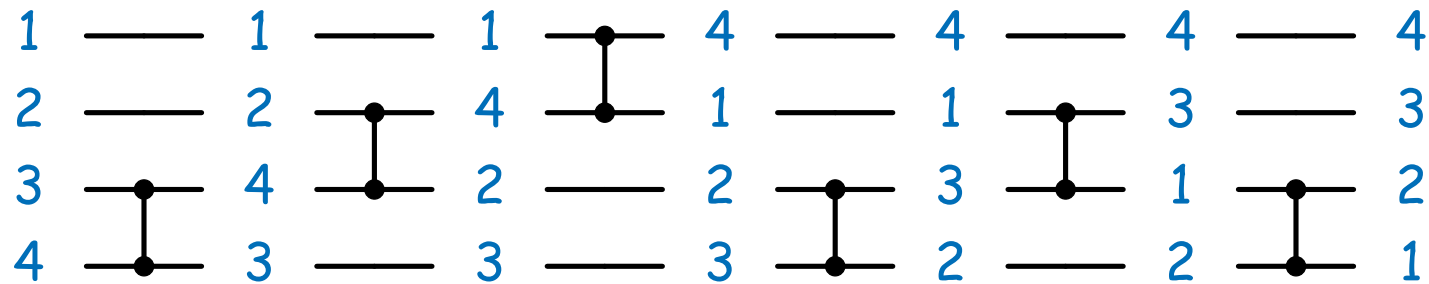
- Sorting a list presents obvious opportunities for parallelization.
- Can illustrate various methods diagrammatically using *comparators* as an elementary unit:



- Each vertical bar represents a *comparator*—a comparison operation or hardware to carry it out—and each horizontal line carries a data item from the list.
- A comparator compares two data items coming from the left, swapping them if the lower one is larger than the upper one.
- Comparators can be grouped into operations that may happen simultaneously; they are always grouped if stacked vertically as in the diagram.

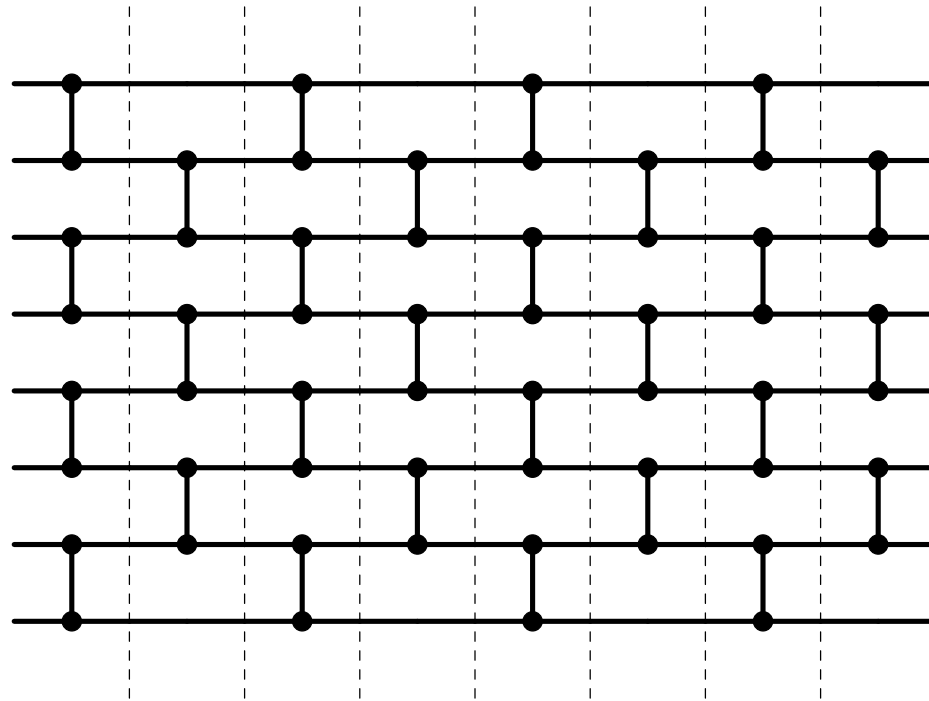
# Sequential sorting

- Here's what a sequential sort (selection sort) might look like:



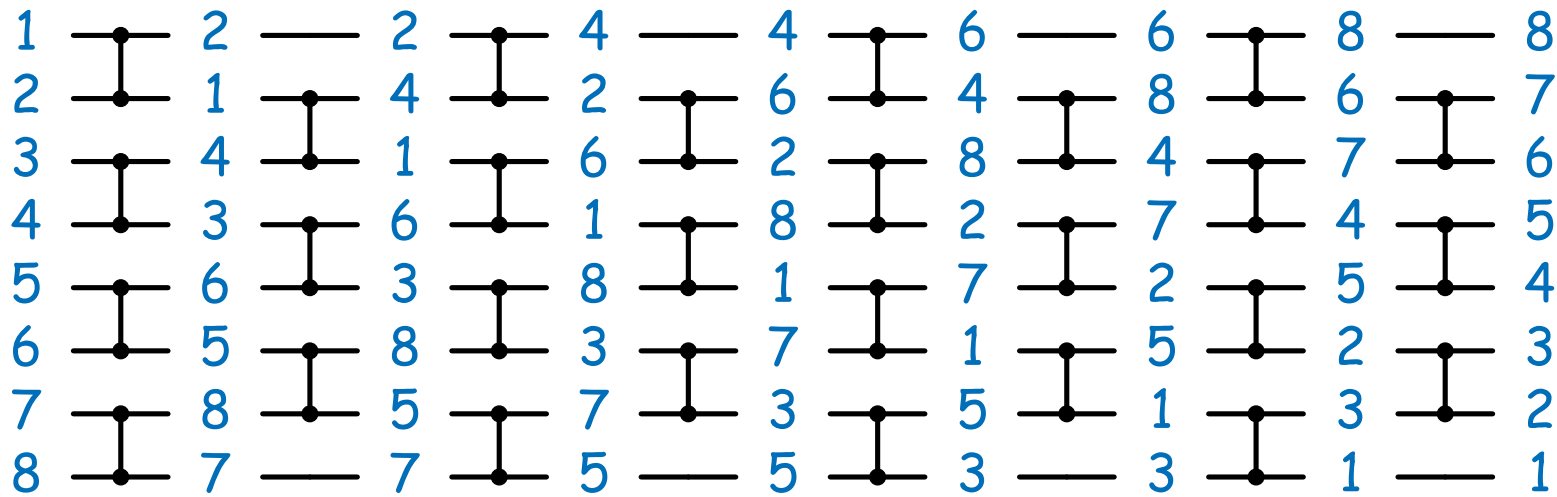
- Each comparator is a separate operation in time.
- In general, there will be  $\Theta(N^2)$  steps.
- But since some comparators operate on distinct data, we ought to be able to overlap operations.

# Odd-Even Transposition Sorter

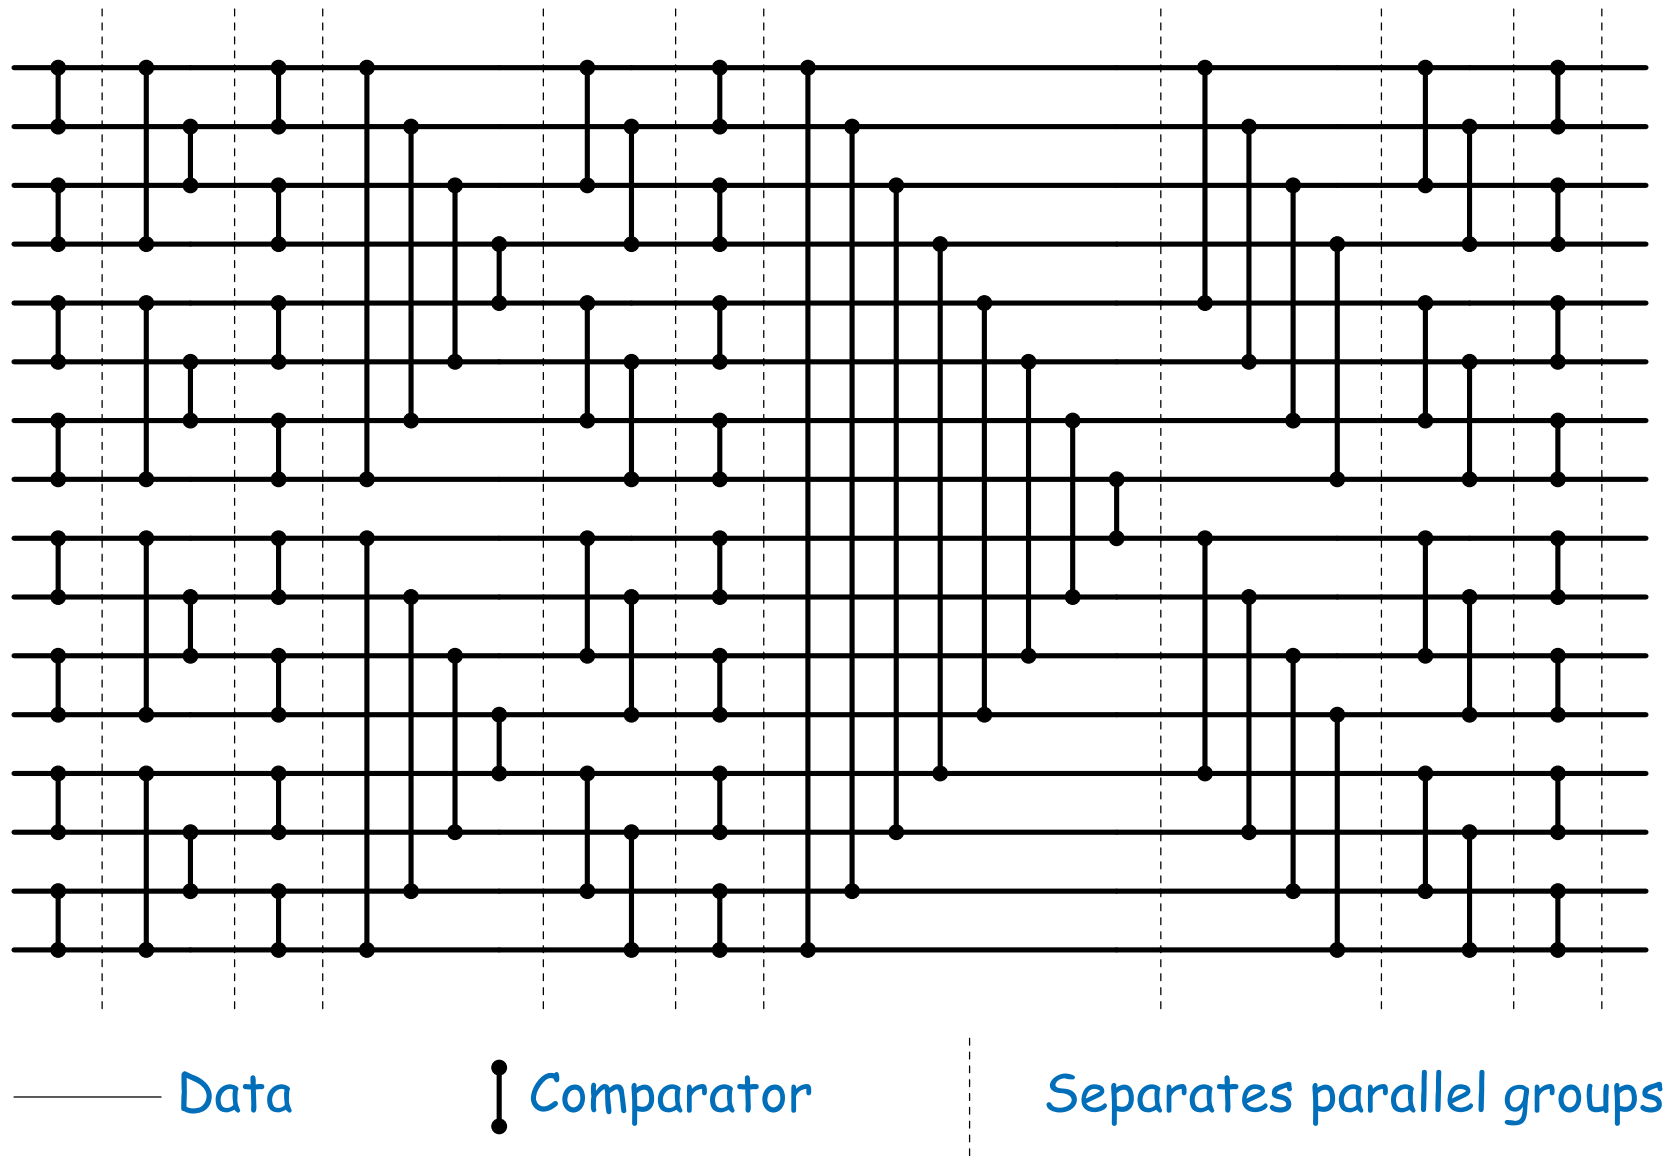


—— Data       Comparator       Separates parallel groups

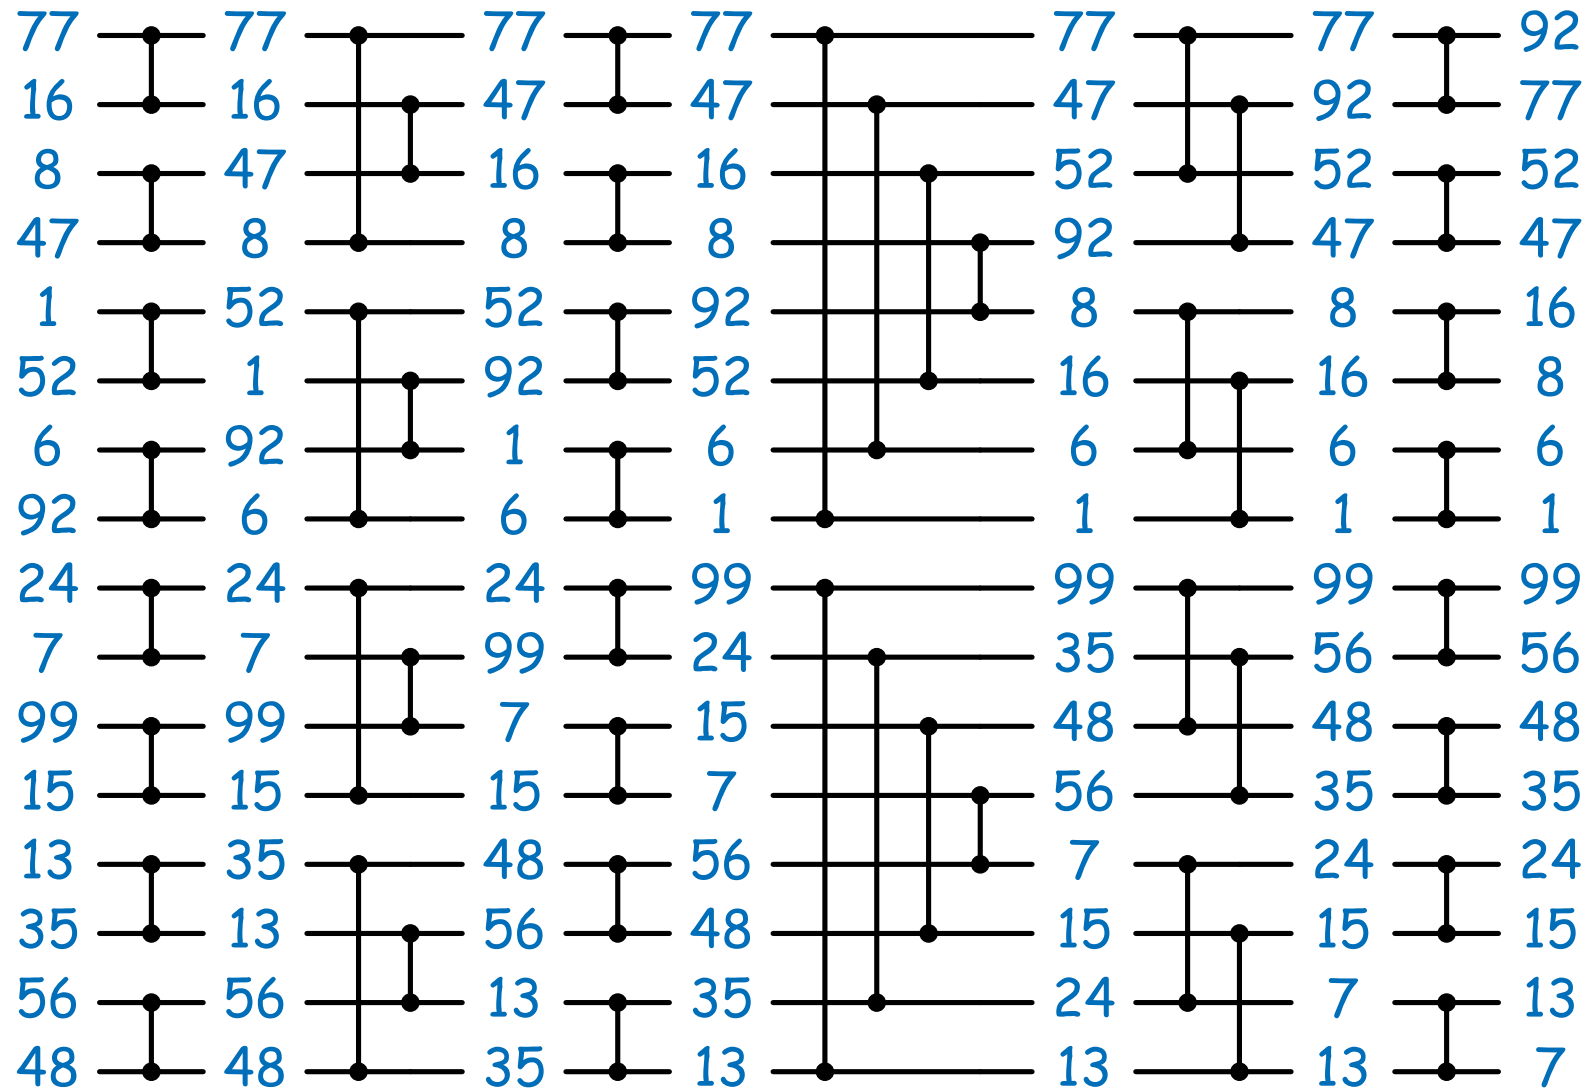
# Odd-Even Sort Example



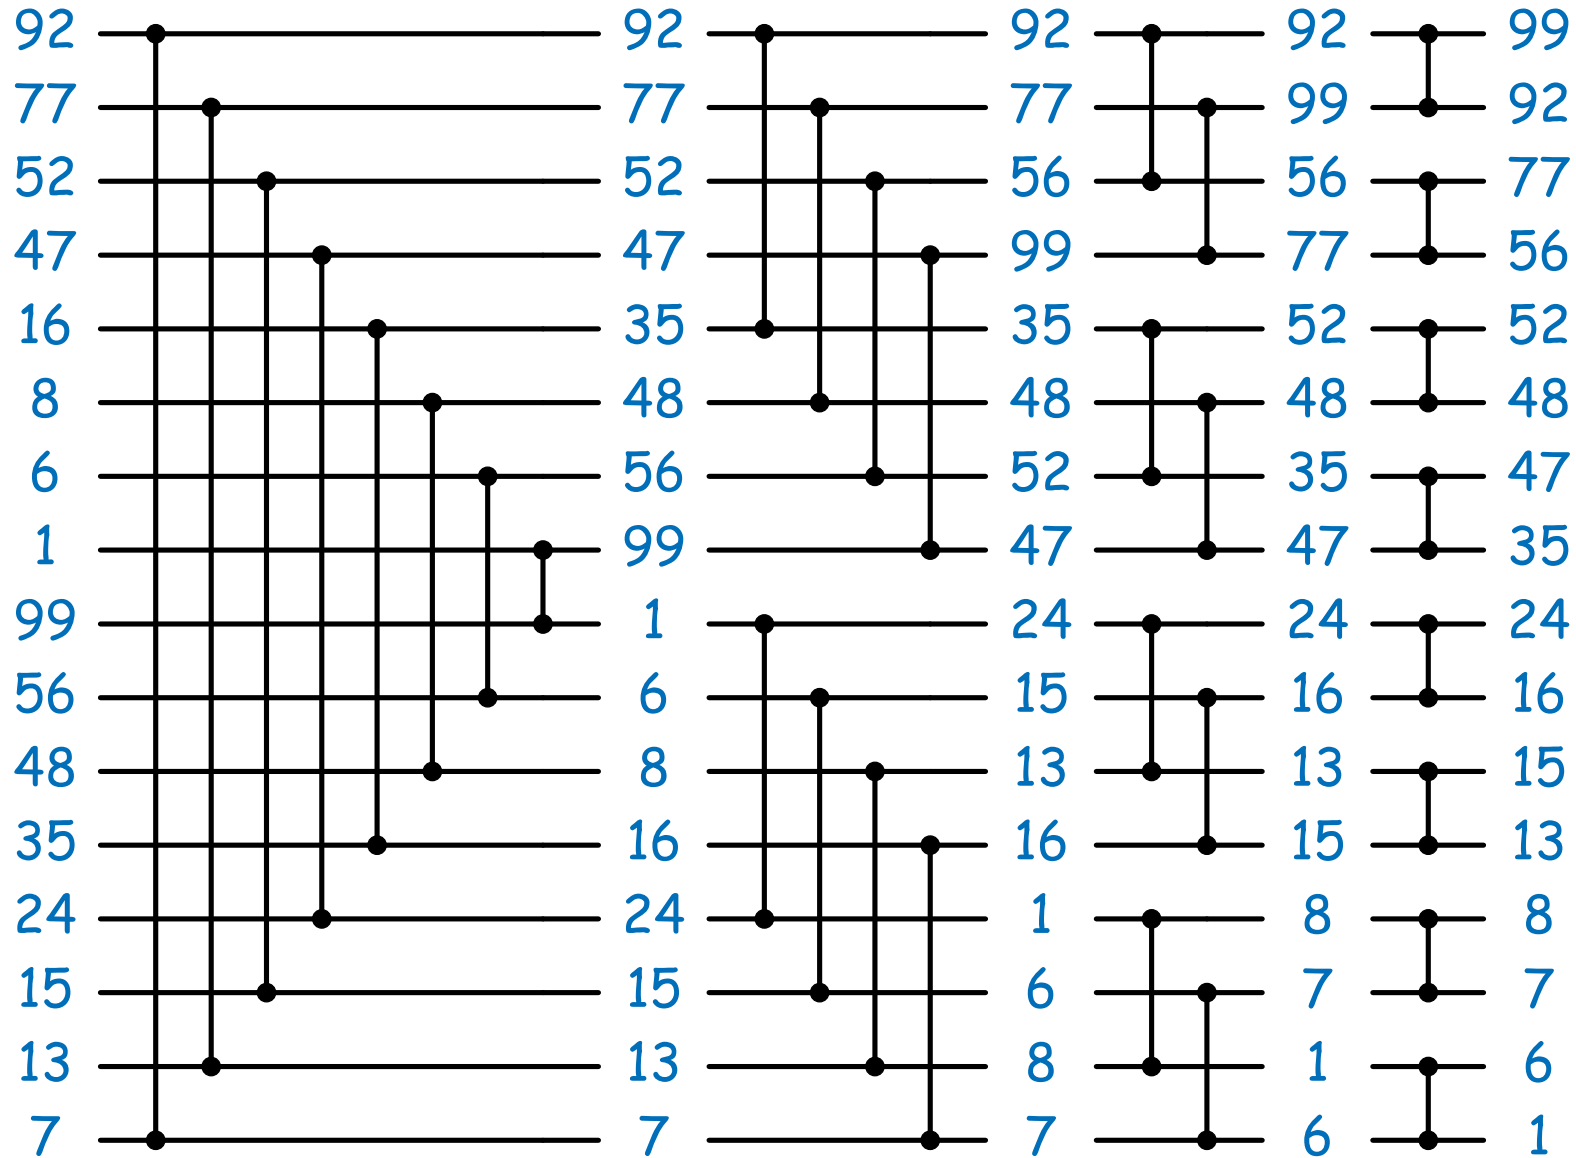
# Example: Bitonic Sorter



# Bitonic Sort Example (I)



# Bitonic Sort Example (II)





# Implementing Parallel Programs

- The sorting diagrams were abstractions.
- Comparators could be processors, or they could be operations divided up among one or more processors.
- Coordinating all of this is the issue.
- One approach is to use *shared memory*, where multiple processors (logical or physical) share one memory.
- This introduces conflicts in the form of *race conditions*: processors racing to access data.

# Memory Conflicts: Abstracting the Essentials

- When considering problems relating to shared-memory conflicts, it is useful to look at the primitive read-to-memory and write-to-memory operations.
- E.g., the program statements on the left cause the actions on the right.

```
x = 5  
x = square(x)
```

```
y = 6  
y += 1
```

```
WRITE 5 -> x  
READ x -> 5  
(calculate 5*5 -> 25)  
WRITE 25 -> x  
WRITE 6 -> y  
READ y -> 6  
(calculate 6+1 -> 7)  
WRITE 7 -> y
```

# Conflict-Free Computation

- Suppose we divide this program into two separate processes, P1 and P2:

```
x = 5  
x = square(x)
```

```
y = 6  
y += 1
```

P1

```
WRITE 5 -> x  
READ x -> 5  
(calculate 5*5 -> 25)  
WRITE 25 -> x
```

P2

```
WRITE 6 -> y  
READ y -> 6  
(calculate 6+1 -> 7)  
WRITE 7 -> y
```

```
x = 25  
y = 7
```

- The result will be the same regardless of which process's READs and WRITEs happen first, because they reference different variables.

# Read-Write Conflicts

- Suppose that both processes read from `x` after it is initialized.

<code>x = 5</code>	
<code>x = square(x)</code>	<code>y = x + 1</code>
P1	P2
<code>READ x -&gt; 5</code> <code>(calculate 5*5 -&gt; 25)</code> <code>WRITE 25 -&gt; x</code> <code> </code>	<code> </code> <code>READ x -&gt; 5</code> <code>(calculate 5+1 -&gt; 6)</code> <code>WRITE 6 -&gt; y</code>
<code>x = 25</code> <code>y = 6</code>	

- The statements in `P2` must appear in the given order, but they need not line up like this with statements in `P1`, because the execution of `P1` and `P2` is independent.

## Read-Write Conflicts (II)

- Here's another possible sequence of events

$x = 5$	
$x = \text{square}(x)$	$y = x + 1$
P1	P2
<pre>READ x -&gt; 5 (calculate 5*5 -&gt; 25) WRITE 25 -&gt; x      </pre>	<pre>      READ x -&gt; 25 (calculate 25+1 -&gt; 26) WRITE 26 -&gt; y</pre>
$x = 25$ $y = 26$	

## Read-Write Conflicts (III)

- The problem here is that nothing forces P1 to wait for P1 to read x before setting it.
- Observation: The “calculate” lines have no effect on the outcome. They represent actions that are entirely local to one processor.
- The effect of “computation” is simply to delay one processor.
- But processors are assumed to be delayable by many factors, such as time-slicing (handing a processor over to another user’s task), or processor speed.
- So the effect of computation adds nothing new to our simple model of shared-memory contention that isn’t already covered by allowing any statement in one process to get delayed by any amount.
- So we’ll just look at READ and WRITE in the future.

# Write-Write Conflicts

- Suppose both processes write to  $x$ :

$x = 5$	
$x = \text{square}(x)$	$x = x + 1$
P1	P2
 READ $x \rightarrow 5$     WRITE 25 $\rightarrow x$	READ $x \rightarrow 5$   WRITE 6 $\rightarrow x$ 
$x = 25$	

- This is a *write-write conflict*: two processes race to be the one that “gets the last word” on the value of  $x$ .

## Write-Write Conflicts (II)

$x = 5$	
$x = \text{square}(x)$	$x = x + 1$
P1	P2
 READ $x \rightarrow 5$ WRITE 25 $\rightarrow x$ 	READ $x \rightarrow 5$     WRITE 6 $\rightarrow x$
$x = 26$	

- This ordering is also possible; P2 gets the last word.
- There are also read-write conflicts here. What is the total number of possible final values for  $x$ ?



## Write-Write Conflicts (II)

$x = 5$	
$x = \text{square}(x)$	$x = x + 1$
P1	P2
 READ $x \rightarrow 5$ WRITE 25 $\rightarrow x$ 	READ $x \rightarrow 5$     WRITE 6 $\rightarrow x$
$x = 26$	

- This ordering is also possible; P2 gets the last word.
- There are also read-write conflicts here. What is the total number of possible final values for  $x$ ? **Four: 25, 5, 26, 36**