Lecture 33: Coordinating Parallel Computation

Let's go back to bank accounts:

```python
class BankAccount:
    def __init__(self, initial_balance):
        self._balance = initial_balance
    @property
    def balance(self): return self._balance
    def withdraw(amount):
        if amount > self._balance:
            raise ValueError("insufficient funds")
        else:
            self._balance -= amount
            return self._balance
```

acct = BankAccount(10)

acct.withdraw(8)	acct.withdraw(7)

- At this point, we'd *like* to have the system raise an exception for one of the two withdrawals, and to set `acct.balance` to either 2 or 3, depending on with withdrawer gets to the bank first, like this...

Desired Outcome

```python
class BankAccount:
    def withdraw(amount):
        if amount > self._balance:
            raise ValueError("insufficient funds")
        else:
            self._balance -= amount
            return self._balance
```

acct = BankAccount(10)

| acct.withdraw(8) | acct.withdraw(7) |

```
READ acct._balance ->  10
WRITE acct._balance -> 2
```

```
                        READ acct._balance -> 2
                        <raise exception>
```

But instead, we might get...

Undesireable Outcome

```
class BankAccount:
    def withdraw(amount):
        if amount > self._balance:
            raise ValueError("insufficient funds")
        else:
            self._balance -= amount
            return self._balance
```

acct = BankAccount(10)

acct.withdraw(8)	acct.withdraw(7)

```
READ acct._balance ->  10        READ acct._balance -> 10
WRITE acct._balance -> 2
                                 WRITE acct._balance -> 3
```

Oops!

Serializability

- We define the desired outcomes as those that would happen if with-drawals happened sequentially, in *some* order.

- The *nondeterminism* as to which order we get is acceptable, but results that are inconsistent with both orderings are not.

- These latter happen when operations overlap, so that the two processes see *inconsistent* views of the account.

- We want the withdrawal operation to act as if it is *atomic*—as if, once started, the operation proceeds without interruption and without any overlapping effects from other operations.

One Solution: Critical Sections

- Some programming languages (e.g., Java) have special syntax for this. In Python, we can arrange something like this:

```python
def withdraw(amount):
    with CriticalSectionManager:
        if amount > self._balance:
            raise ValueError("insufficient funds")
        else:
            self._balance -= amount
            return self._balance
```

- The `with` construct:

  1. Calls the `__enter__()` method of its "context manager" argument (here, some object we'll call `CriticalSectionManager`);
  2. Executes the body (indented portion);
  3. Finally, it calls the `__exit__()` method on the context manager. It guarantees that it will *always* do so, no matter how you exit from the body (via `return`, exception, etc.).

- The idea is that our *CriticalSectionManager* object should let just one process through at a time. How?

Locks

- To implement our critical sections, we'll need some help from the operating system or underlying hardware.

- A common low-level construct is the *lock* or *mutex* (for "mutual exclusion"): an object that at any given time is "owned" by one process.

- If `L` is a lock, then

  - `L.acquire()` attempts to own `L` on behalf of the calling process. If someone else owns it, the caller *waits* for it to be release.
  - `L.release()` relinquishes ownership of `L` (if the calling process owns it).

Implementing Critical Regions

- Using locks, it's easy to create the desired context manager:

```
from threading import Lock

class CriticalSection:
    def __init__(self):
        self.__lock = Lock()

    def __enter__(self):
        self.__lock.acquire()

    def __exit__(self, exception_type, exception_val, traceback):
        self.__lock.release()
CriticalSectionManager = CriticalSection()
```

- The extra arguments to `__exit__` provide information about the exception, if any, that caused the **with** body to be exited.

- (In fact, the bare `Lock` type itself already has `__enter__` and `__exit__` procedures, so you don't really have to define an extra type).

Granularity

- We've envisioned critical sections as being atomic with respect to *all* other critical sections.

- Has the advantage of simplicity and safety, but causes unnecessary waits.

- In fact, different accounts need not coordinate with each other. We can have a separate critical section manager (or lock) for each account object:

```python
class BankAccount:
    def __init__(self, initial_balance):
        self._balance = initial_balance
        self._critical = CriticalSection()
    def withdraw(self, amount):
        with self._critical:
            ...
```

- That is, can produce a solution with finer *granularity* of locks.

Synchronization

- Another kind of problem arises when different processes must communicate. In that case, one may have to wait for the other to send something.

- This, for example, doesn't work too well:

```python
class Mailbox:
    def __init__(self):
        self._queue = []
    def deposit(self, msg):
        self._queue.append(msg)
    def pickup(self):
        while not self._queue:
            pass
        return self._queue.pop()
```

- Idea is that one process deposits a message for another to pick up later.

- What goes wrong?

Problems with the Naive Mailbox

```
class Mailbox:
    def __init__(self):
        self._queue = []
    def deposit(self, msg):
        self._queue.append(msg)
    def pickup(self):
        while not self._queue:
            pass
        return self._queue.pop()
```

- *Inconsistency:* Two processes picking up mail can find the queue occupied simultaneously, but only one will succeed in picking up mail, and the other will get exception.

- *Busy-waiting:* The loop that waits for a message uses up processor time.

- *Deadlock:* If one is running two logical processes on one processor, busy-waiting can lead to nobody making any progress.

- *Starvation:* Even without busy-waiting one process can be shut out from ever getting mail.

Conditions

- One way to deal with this is to augment locks with *conditions:*

```python
from threading import Condition
class Mailbox:
    def __init__(self):
        self._queue = []
        self._condition = Condition()
    def deposit(self, msg):
        with self._condition:
            self._queue.append(msg)
            self._condition.notify()
    def pickup(self):
        with self._condition:
            while not self._queue:
                self._condition.wait()
            return self._queue.pop()
```

- Conditions act like locks with methods `wait`, `notify` (and others).

- `wait` releases the lock, waits for someone to call `notify`, and then reacquires the lock.

Another Approach: Messages

- Turn the problem inside out: instead of client processes deciding how to coordinate their operations on data, let the *data* coordinate its actions.

- From the Mailbox's perspective, things look like this:

```
self._queue = []
while True:
    wait for a request, R, to deposit or pickup
    if R is a deposit of msg:
        self.__queue.append(msg)
        send back acknowledgement
    elif self.__queue and R is a pickup:
        msg = self.__queue.pop()
        send back msg
```

Rendezvous

- Following ideas from C.A.R Hoare, the Ada language used the notion of a *rendezvous* for this purpose:

```
task type Mailbox is
    entry deposit(Msg: String);
    entry pickup(Msg: out String);
end Mailbox;

task body Mailbox is
    Queue: ...
begin
    loop
        select
          accept deposit(Msg: String) do Queue.append(Msg); end;
        or when not Queue.empty =>
          accept pickup(Msg: out String) do Queue.pop(Msg); end;
        end select;
    end loop;
end;
```